

## Assignment 3

### GitHub Repo:

<https://github.com/xq443/distributed-system/tree/main/Assignment/3>

### Redis DB Design:

#### 1. Redis Key Design

##### Examples

- Skier Data
  - o Key: skier:{skierID}:season:{seasonID}
  - o Value: Hash with fields like:
    - days\_skied (Set of dayIDs skied).
    - vertical:{dayID} (Vertical total for the day).
    - lifts:{dayID} (List of lifts ridden).
- Resort Visitors
  - o Key: resort:{resortID}:day:{dayID}:visitors
  - o Value: Set of skierIDs.

#### 2. API Retrieval Logic

##### API Endpoint:

GET /skier/{skierID}/season/{seasonID}/summary

##### Steps for Data Retrieval:

1. Retrieve Days Skied  
Query: SCARD skier:{skierID}:season:{seasonID}:days\_skied
  - o Efficiently counts unique skiing days.
2. Retrieve Vertical Totals  
Query: HGETALL skier:{skierID}:season:{seasonID}:vertical
  - o Returns all day-to-vertical mappings for the skier.
3. Retrieve Lifts by Day  
For each dayID: Query: LRANGE skier:{skierID}:season:{seasonID}:lifts:{dayID} 0 -1
  - o Returns the list of lifts ridden on a specific day.
4. Response Construction  
Combine the results into a JSON response summarizing days skied, total vertical feet, and lifts ridden per day.

#### 3. Optimal Key Design Trade-offs

##### Chosen Key Structure

- Example: skier:{skierID}:season:{seasonID}:vertical
  - o Pros:
    - Optimized for skier-centric queries (e.g., vertical totals per day).
    - Prevents unnecessary data fetching by scoping to the skier and season.
  - o Cons:
    - Separate keys for vertical and lifts can require multiple queries for full retrieval.

##### Alternative Design

- Key: skier:{skierID}:season:{seasonID}:day:{dayID}
  - Pros:
    - Each day's data (vertical and lifts) is stored together for simpler retrieval.
    - Fewer Redis keys overall.
  - Cons:
    - Requires fetching and aggregating all day:{dayID} keys for season-level queries.
    - Increases query complexity for aggregating totals.

#### 4. Trade-offs between Redis and other database choices

Redis is less suited for complex queries or durability but excels in low-latency, high-throughput use cases. If real-time performance is the priority, Redis is the optimal choice, while DynamoDB or MySQL would be better for large-scale durability or relational needs.

#### Deployment Topology on AWS:

The deployment is composed of an EC2 instance and redis within AWS.

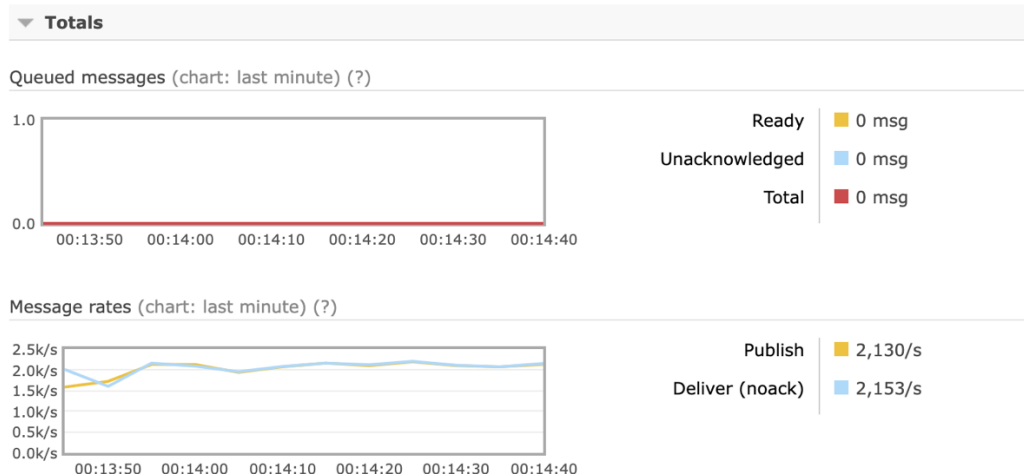
Redis resides in the same EC2 instance with consumer application that reads messages from RabbitMQ; Tomcat server and RabbitMQ resides in another EC2 instance correspondingly.

#### Throughput Performance:

Instance type

z1d.large \* 3

### Overview



Thread counts: 290

Successful requests: 200000

Failed requests: 0

Total time: 89765 ms

Throughput: 2228.0398819138863 requests/second